

# Final Project Submission

Please fill out:

- Student name:
- Student pace: self paced / part time / full time
- Scheduled project review date/time:
- Instructor name:
- Blog post URL:

```
In [1]: 1  #Importing Required Libraries
        2
        3  import numpy as np
        4  import pandas as pd
        5
        6  from sklearn.tree import DecisionTreeClassifier
        7  from sklearn.ensemble import RandomForestClassifier
        8  from sklearn.tree import DecisionTreeRegressor
        9  from sklearn.neighbors import KNeighborsClassifier
       10  from sklearn.linear_model import LogisticRegression
       11  from sklearn.svm import SVC
       12  from sklearn.metrics import precision_score, recall_score, accuracy_score
       13
       14
       15  import matplotlib.pyplot as plt
       16
```

```
In [2]: 1  #Create initial DF
        2
        3  df = pd.read_csv('ChurnData.csv')
```

```
In [3]: 1 #Checking Info
        2
        3 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   state                                3333 non-null   object
1   account length                       3333 non-null   int64
2   area code                           3333 non-null   int64
3   phone number                         3333 non-null   object
4   international plan                   3333 non-null   object
5   voice mail plan                      3333 non-null   object
6   number vmail messages                3333 non-null   int64
7   total day minutes                    3333 non-null   float64
8   total day calls                      3333 non-null   int64
9   total day charge                     3333 non-null   float64
10  total eve minutes                     3333 non-null   float64
11  total eve calls                      3333 non-null   int64
12  total eve charge                     3333 non-null   float64
13  total night minutes                  3333 non-null   float64
14  total night calls                    3333 non-null   int64
15  total night charge                   3333 non-null   float64
16  total intl minutes                   3333 non-null   float64
17  total intl calls                     3333 non-null   int64
18  total intl charge                    3333 non-null   float64
19  customer service calls                3333 non-null   int64
20  churn                                3333 non-null   bool
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB
```

```
In [4]: 1 #Get data set dimensions
        2
        3 df.shape
```

```
Out[4]: (3333, 21)
```

```
In [5]: 1 #Get data set
        2
        3 df.describe()
```

Out[5]:

	account length	area code	number vmail messages	total day minutes	total day calls	total day charge	total eve minutes
<b>count</b>	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000
<b>mean</b>	101.064806	437.182418	8.099010	179.775098	100.435644	30.562307	200.980348
<b>std</b>	39.822106	42.371290	13.688365	54.467389	20.069084	9.259435	50.713844
<b>min</b>	1.000000	408.000000	0.000000	0.000000	0.000000	0.000000	0.000000
<b>25%</b>	74.000000	408.000000	0.000000	143.700000	87.000000	24.430000	166.600000
<b>50%</b>	101.000000	415.000000	0.000000	179.400000	101.000000	30.500000	201.400000
<b>75%</b>	127.000000	510.000000	20.000000	216.400000	114.000000	36.790000	235.300000
<b>max</b>	243.000000	510.000000	51.000000	350.800000	165.000000	59.640000	363.700000

```
In [6]: 1 #Familiarizing with Target Variable
        2
        3 df['churn'].value_counts()
```

Out[6]: False 2850  
True 483  
Name: churn, dtype: int64

```
In [7]: 1 #Further investigation of target variable
        2
        3 df['churn'].value_counts(normalize = True)
```

Out[7]: False 0.855086  
True 0.144914  
Name: churn, dtype: float64

```
In [8]: 1 from sklearn.model_selection import train_test_split
        2
        3 #Define X & Y
        4
        5 X = df.drop('churn', axis = 1)
        6 y = df['churn']
        7
        8 #Train Test Split
        9
        10 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1)
```

```

In [9]: 1#Get numeric cols, categorical cols, and cols to frequency
2
3#taken from: https://github.com/flatiron-school/Online-DS-PT-022221/blob/
4
5num_cols = []
6cols_to_ohe = []
7cols_to_freq = []
8
9for c in X_train.columns:
10     # Want to grab numeric columns
11     if X_train[c].dtype in ['float64', 'int64']:
12         # same as if X_train[c].dtype == 'float64'
13         num_cols.append(c)
14
15     # Then grab columns with fewer than 10 unique values
16     elif len(X_train[c].unique()) < 10:
17         cols_to_ohe.append(c)
18
19     # Then grab columns with more than 10, since we won't OHE those
20     else:
21         cols_to_freq.append(c)
22
23print(f'Numeric Columns: ')
24print(num_cols)
25print(f'==')
26
27print(f'OHE Columns: ')
28print(cols_to_ohe)
29print(f'==')
30
31print(f'Categorical Columns that will not OHE: ')
32print(cols_to_freq)

```

Numeric Columns:

```
['account length', 'area code', 'number vmail messages', 'total day minutes', 'total day calls', 'total day charge', 'total eve minutes', 'total eve calls', 'total eve charge', 'total night minutes', 'total night calls', 'total night charge', 'total intl minutes', 'total intl calls', 'total intl charge', 'customer service calls']
```

==

OHE Columns:

```
['international plan', 'voice mail plan']
```

==

Categorical Columns that will not OHE:

```
['state', 'phone number']
```

In [10]: *Creating Pre Processing Pipeline*

```

2
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import MinMaxScaler, OneHotEncoder
import category_encoders as ce

8
9
Sourced From: https://github.com/flatiron-school/Online-DS-PT-022221/blob/main
11
Numerical Pre Processing Steps:
13
num_transformer = Pipeline(steps=[
15('num_imputer', SimpleImputer(strategy='mean')),
16('scaler', MinMaxScaler())])
17
One Hot Pre processing Steps:
19
ohe_transformer = Pipeline(steps=[
21('ohe_imputer', SimpleImputer(strategy='constant', fill_value='Unknown')),
22('ohencoder', OneHotEncoder(handle_unknown='ignore'))])
23
Frequency Columns Pre processing Steps:
25
freq_transformer = Pipeline(steps=[
27('freq_imputer', SimpleImputer(strategy='constant', fill_value='Unknown')),
28('freq_enc', ce.count.CountEncoder(normalize=True,
29                                     handle_unknown=0,
30                                     min_group_size=0.001,
31                                     min_group_name='Other'))])
32
Put all Processing Steps together in Columntransformer
34
preprocessor = ColumnTransformer(
36transformers=[
37    ('num', num_transformer, num_cols),
38    ('ohe', ohe_transformer, cols_to_ohe),
39    ('freq', freq_transformer, cols_to_freq)])

```

In [11]: *#Fit preprocessor to X\_Train data*

```

2
3 X_train_proc = pd.DataFrame(preprocessor.fit_transform(X_train))

```

```

In [12]: 1 #Get feature names from processor
2
3 #Sourced from: https://johaupt.github.io/scikit-learn/tutorial/python/d
4
5 import warnings
6 import sklearn
7 import pandas as pd
8
9 def get_feature_names(column_transformer):
10     """Get feature names from all transformers.
11     Returns
12     -----
13     feature_names : list of strings
14         Names of the features produced by transform.
15     """
16     # Remove the internal helper function
17     #check_is_fitted(column_transformer)
18
19     # Turn loopkup into function for better handling with pipeline late
20     def get_names(trans):
21         # >> Original get_feature_names() method
22         if trans == 'drop' or (
23             hasattr(column, '__len__') and not len(column)):
24             return []
25         if trans == 'passthrough':
26             if hasattr(column_transformer, '_df_columns'):
27                 if ((not isinstance(column, slice))
28                     and all(isinstance(col, str) for col in column)):
29                     return column
30             else:
31                 return column_transformer._df_columns[column]
32             else:
33                 indices = np.arange(column_transformer._n_features)
34                 return ['x%d' % i for i in indices[column]]
35         if not hasattr(trans, 'get_feature_names'):
36             # >>> Change: Return input column names if no method available
37             # Turn error into a warning
38             warnings.warn("Transformer %s (type %s) does not "
39                           "provide get_feature_names. "
40                           "Will return input column names if ava
41                           % (str(name), type(trans).__name__))
42             # For transformers without a get_features_names method, use
43             # names to the column transformer
44             if column is None:
45                 return []
46             else:
47                 return [name + "__" + f for f in column]
48
49         return [name + "__" + f for f in trans.get_feature_names()]
50
51     ### Start of processing
52     feature_names = []
53
54     # Allow transformers to be pipelines. Pipeline steps are named diff
55     if type(column_transformer) == sklearn.pipeline.Pipeline:
56         l_transformers = [(name, trans, None, None) for step, name, tra

```

```

57     else:
58         # For column transformers, follow the original method
59         l_transformers = list(column_transformer._iter(fitted=True))
60
61
62     for name, trans, column, _ in l_transformers:
63         if type(trans) == sklearn.pipeline.Pipeline:
64             # Recursive call on pipeline
65             _names = get_feature_names(trans)
66             # if pipeline has no transformer that returns names
67             if len(_names)==0:
68                 _names = [name + "__" + f for f in column]
69                 feature_names.extend(_names)
70         else:
71             feature_names.extend(get_names(trans))
72
73     return feature_names

```

```

In [13]: 1 #Get feature names of preprocessor
2
3 processed_feature_names = get_feature_names(preprocessor)
4
5 #Assign column names to preprocessed X_train_cols
6
7 X_train_proc.columns = processed_feature_names
8
9 #transform X_test
10
11 X_test_proc = pd.DataFrame(preprocessor.transform(X_test))

```

```

<ipython-input-12-0d499a1f2f22>:38: UserWarning: Transformer num_imputer
(type SimpleImputer) does not provide get_feature_names. Will return input
column names if available
    warnings.warn("Transformer %s (type %s) does not "
<ipython-input-12-0d499a1f2f22>:38: UserWarning: Transformer scaler (type
MinMaxScaler) does not provide get_feature_names. Will return input column
names if available
    warnings.warn("Transformer %s (type %s) does not "
<ipython-input-12-0d499a1f2f22>:38: UserWarning: Transformer ohe_imputer
(type SimpleImputer) does not provide get_feature_names. Will return input
column names if available
    warnings.warn("Transformer %s (type %s) does not "
<ipython-input-12-0d499a1f2f22>:38: UserWarning: Transformer freq_imputer
(type SimpleImputer) does not provide get_feature_names. Will return input
column names if available
    warnings.warn("Transformer %s (type %s) does not "
<ipython-input-12-0d499a1f2f22>:38: UserWarning: Transformer freq_enc (type
CountEncoder) does not provide get_feature_names. Will return input column
names if available
    warnings.warn("Transformer %s (type %s) does not "

```

```

In [14]: 1 #Get Metrics for Classifier Model Types:
2
3 from sklearn.tree import DecisionTreeClassifier
4 from sklearn.ensemble import RandomForestClassifier
5 from sklearn.tree import DecisionTreeRegressor
6 from sklearn.neighbors import KNeighborsClassifier
7 from sklearn.linear_model import LogisticRegression
8 from sklearn.svm import SVC
9
10 def multi_model_analysis(models, weight):
11
12     classifier = []
13     precision = []
14     recall = []
15     accuracy = []
16     f1 = []
17
18     # Class Weight Models
19
20     def get_weighted_model_metrics(reg, weight, X_train_proc, y_train,
21
22         from sklearn.metrics import precision_score, recall_score, accu
23
24         model = reg(class_weight = weight)
25
26         model.fit(X_train_proc, y_train)
27
28         model_pred_train = model.predict(X_train_proc)
29
30         model_pred_test = model.predict(X_test_proc)
31
32         #Get Evaluation Metrics for DTC Train Set:
33
34         model_train_precision = precision_score(model_pred_train, y_train)
35         model_train_recall = recall_score(model_pred_train, y_train)
36         model_train_accuracy = accuracy_score(model_pred_train, y_train)
37         model_train_f1 = f1_score(model_pred_train, y_train)
38
39         #Get Evaluation Metrics for DTC Test Set:
40
41         model_test_precision = precision_score(model_pred_test, y_test)
42         model_test_recall = recall_score(model_pred_test, y_test)
43         model_test_accuracy = accuracy_score(model_pred_test, y_test)
44         model_test_f1 = f1_score(model_pred_test, y_test)
45
46         classifier.append(reg)
47         precision.append(model_test_precision)
48         recall.append(model_test_recall)
49         accuracy.append(model_test_accuracy)
50         f1.append(model_test_f1)
51
52     #Non Class Weight Models
53
54     def get_model_metrics(reg, X_train_proc, y_train, X_test_proc, y_test)
55
56         from sklearn.metrics import precision_score, recall_score, accu

```



```


57
58     model = reg()
59
60     model.fit(X_train_proc, y_train)
61
62     model_pred_train = model.predict(X_train_proc)
63
64     model_pred_test = model.predict(X_test_proc)
65
66     #Get Evaluation Metrics for DTC Train Set:
67
68     model_train_precision = precision_score(model_pred_train, y_train)
69     model_train_recall = recall_score(model_pred_train, y_train)
70     model_train_accuracy = accuracy_score(model_pred_train, y_train)
71     model_train_f1 = f1_score(model_pred_train, y_train)
72
73     #Get Evaluation Metrics for DTC Test Set:
74
75     model_test_precision = precision_score(model_pred_test, y_test)
76     model_test_recall = recall_score(model_pred_test, y_test)
77     model_test_accuracy = accuracy_score(model_pred_test, y_test)
78     model_test_f1 = f1_score(model_pred_test, y_test)
79
80     classifier.append(reg)
81     precision.append(model_test_precision)
82     recall.append(model_test_recall)
83     accuracy.append(model_test_accuracy)
84     f1.append(model_test_f1)
85
86     for item in models:
87         if item in [DecisionTreeClassifier, RandomForestClassifier, LogisticRegression]:
88             get_weighted_model_metrics(item, weight, X_train_proc, y_train, X_test_proc)
89         else:
90             get_model_metrics(item, X_train_proc, y_train, X_test_proc)
91
92     metrics_df = pd.concat([pd.DataFrame(classifier), pd.DataFrame(precision),
93                             pd.DataFrame(recall), pd.DataFrame(accuracy),
94                             pd.DataFrame(f1)], axis = 1)
95
96     metrics_df.columns = ['Classifier Model', 'Precision', 'Recall', 'Accuracy', 'F1']
97
98     metrics_df.sort_values(by = 'F1', inplace = True, ascending = True)
99
100    metrics_df.plot(x = 'Classifier Model', y=['Precision', 'Recall', 'Accuracy', 'F1'],
101                    figsize = (15, 10))
102
103    plt.axvline(0.8, 0, 1, color = 'red')
104    plt.title('Model Analysis Ranked by F1')
105
106    return metrics_df

```

```
In [15]: 1 #Test All Models
2
3 models = [DecisionTreeClassifier, RandomForestClassifier, LogisticRegression,
4           KNeighborsClassifier]
5
6 multi_model_analysis(models, None)
```

Out[15]:

	Classifier Model	Precision	Recall	Accuracy	F1
2	<class 'sklearn.linear_model._logistic.LogisticRegression'>	0.166667	0.590909	0.852	0.260000
3	<class 'sklearn.svm._classes.SVC'>	0.333333	0.838710	0.886	0.477064
5	<class 'sklearn.neighbors._classification.KNeighborsClassifier'>	0.346154	0.843750	0.888	0.490909
0	<class 'sklearn.tree._classes.DecisionTreeClassifier'>	0.756410	0.662921	0.902	0.706587
4	<class 'sklearn.tree._classes.DecisionTreeRegressor'>	0.756410	0.662921	0.902	0.706587
1	<class 'sklearn.ensemble._forest.RandomForestClassifier'>	0.730769	0.966102	0.954	0.832117



## Initial Multi Model Findings

1. The highest F1 scores come from Decision Tree Classifiers and Decision Tree Regression
2. Random Forest Metrics are showing very promising Accuracy and Recall, however Precision is very low
3. Scores may be improved across all models by first solving for the Class Imbalance


## Check Class Weight Imbalance

1. Multi Model Check will be run with Class Weight Imbalanced rectified

```
In [16]: 1 multi_model_analysis(models, 'balanced')
```

Out[16]:

	Classifier Model	Precision	Recall	Accuracy	F1
5	<class 'sklearn.neighbors._classification.KNei...>	0.346154	0.843750	0.888	0.490909
2	<class 'sklearn.linear_model._logistic.Logisti...>	0.756410	0.380645	0.770	0.506438
3	<class 'sklearn.svm._classes.SVC'>	0.884615	0.570248	0.878	0.693467
4	<class 'sklearn.tree._classes.DecisionTreeRegr...>	0.743590	0.690476	0.908	0.716049
0	<class 'sklearn.tree._classes.DecisionTreeClas...>	0.794872	0.784810	0.934	0.789809
1	<class 'sklearn.ensemble._forest.RandomForestC...>	0.717949	0.982456	0.954	0.829630



## Class Weight Analysis

1. Random Forests yielded the greatest improvement in total F1
2. Decision Tree Classifiers came in a close second
3. Logistic Regression maintained a relatively low score

## Next Steps:

1. Run CV/Grid Search on Random Forest Classifier
2. Determine feature importance

```
In [17]: 1 # #Grid Search Parameters
2
3 # # dt_param_grid = {
4 # #     'criterion': ['gini', 'entropy'],
5 # #     'max_depth': [None, 2, 3, 4, 5, 6],
6 # #     'min_samples_split': [2, 5, 10],
7 # #     'min_samples_leaf': [1, 2, 3, 4, 5, 6]
8 # # }
9
10 # dt_param_grid = {
11 #     'criterion': ['gini', 'entropy'],
12 #     'max_depth': [None, 2, 3, 4, 5, 6],
13 #     'min_samples_split': [2, 5, 10],
14 #     'min_samples_leaf': [1, 2, 3, 4, 5, 6]
15 # }
16
17 # #Instantiate Grid Search CV
18
19 # from sklearn.model_selection import GridSearchCV
20
21 # grid = GridSearchCV(model, dt_param_grid, cv = 3, scoring = 'f1')
22
23 # grid.fit(X_train_proc, y_train)
24
25
```

```
In [18]: 1 # #Evaluating Metrics from Gridsearch CV:
2
3 # print(f'Best Score: {grid.best_score_}')
4
5 # #Evaluating best
6
7 # print(f'Best Parameters: {grid.best_params_}')
8
9 print(f'Best Metrics: criterion: entropy, max_depth: None, min_samples_
```

Best Metrics: criterion: entropy, max\_depth: None, min\_samples\_leaf: 5, min\_samples\_split: 2

```

In [19]: 1 model = RandomForestClassifier(class_weight = 'balanced',
2                                     criterion = 'entropy',
3                                     max_depth = None,
4                                     min_samples_leaf = 5,
5                                     min_samples_split = 2)
6
7 model.fit(X_train_proc, y_train)
8
9 model_pred_train = model.predict(X_train_proc)
10
11 model_pred_test = model.predict(X_test_proc)
12
13 #Get Evaluation Metrics for DTC Train Set:
14
15 model_train_precision = precision_score(model_pred_train, y_train)
16 model_train_recall = recall_score(model_pred_train, y_train)
17 model_train_accuracy = accuracy_score(model_pred_train, y_train)
18 model_train_f1 = f1_score(model_pred_train, y_train)
19
20 #Get Evaluation Metrics for DTC Test Set:
21
22 model_test_precision = precision_score(model_pred_test, y_test)
23 model_test_recall = recall_score(model_pred_test, y_test)
24 model_test_accuracy = accuracy_score(model_pred_test, y_test)
25 model_test_f1 = f1_score(model_pred_test, y_test)
26
27 print(f'Random Forest Model Metrics:')
28 print(f'')
29 print(f'Model Test Precision: {model_test_precision}')
30 print(f'Model Test Recall: {model_test_recall}')
31 print(f'Model Test Accuracy: {model_test_accuracy}')
32 print(f'Model Test F1: {model_test_f1}')
33 print(f'---')
34 print(f'---')
35
36 metrics = ['Precision', 'Recall', 'Accuracy', 'F1']
37 scores = [model_test_precision, model_test_recall, model_test_accuracy,
38           model_test_f1]
39 plt.barh(metrics, scores)
40 plt.title('Updated Random Forest Metrics with Grid Search')
41 plt.axvline(0.8, 0, 1, color = 'red')
42
43

```

Random Forest Model Metrics:

Model Test Precision: 0.8333333333333334

Model Test Recall: 0.8666666666666667


Model Test Accuracy: 0.954

Model Test F1: 0.8496732026143791

---

---

Out[19]: <matplotlib.lines.Line2D at 0x7fc28893ee20>



## Grid Search CV Results

Overall scoring has gone up significantly, with Accuracy hitting 95% and F1 increasing to 85%

## Next Steps


1. Determine Feature Importance
2. Establish final steps of qualitative analysis

```

In [20]: 1 #Establish Feature Importance in a Dta Frame
          2
          3 feature_by_importance = pd.concat([pd.DataFrame(X_train_proc.columns),
          4                                         axis = 1)
          5
          6 feature_by_importance.columns = ['Feature', 'Importance']
          7
          8 feature_by_importance.sort_values(by='Importance', ascending = True, in
          9
          10 plot = plt.figure()
          11 plot.set_figwidth(4)
          12 plot.set_figheight(10)
          13
          14 plt.title('Feature Importances')
          15
          16 plt.barh(feature_by_importance['Feature'],
          17          feature_by_importance['Importance'])

```

Out[20]: <BarContainer object of 22 artists>



## Feature Importance Results

1. Top features include Customer Service Calls, Total Day Minutes, Total Day Charge, Total Night Minutes, Total Night Charge
2. It appears that the most active, high frequency users are most likely to churn out

## Next Steps

1. Establish average usage and point in which users are most likely to churn



```
In [21]: 1 #Service Calls Analysis
2
3 cust_service_calls_mean = np.mean(df['customer service calls'])
4
5 cust_service_calls_sd = np.std(df['customer service calls'])
6
7 non_churn_cust_service_calls_mean = np.mean(df[df['churn'] == 0]['custo
8
9 churn_cust_service_calls_mean = np.mean(df[df['churn'] == 1]['customer
10
11 print(f'Mean of Customer Service Calls For Non Churners: {non_churn_cus
12 print(f'Standard Deviation of Customer Service Calls: {cust_service_cal
13 print(f'Mean of Customer Service Calls for Churners: {churn_cust_servic
```

Mean of Customer Service Calls For Non Churners: 1.4498245614035088

Standard Deviation of Customer Service Calls: 1.3152936866709521

Mean of Customer Service Calls for Churners: 2.229813664596273

```
In [22]: 1 #Day Time Minutes Analysis
2
3 day_miniutes_mean = np.mean(df['total day minutes'])
4
5 day_minutes_sd = np.std(df['total day minutes'])
6
7 non_churn_day_minutes_mean = np.mean(df[df['churn'] == 0]['total day mi
8
9 churn_day_minutes_mean = np.mean(df[df['churn'] == 1]['total day minute
10
11 print(f'Mean of total day minutes Non Churners: {non_churn_day_minutes_
12 print(f'Standard Deviation of total day minutes: {day_minutes_sd}')
13 print(f'Mean of total day minutes for Churners: {churn_day_minutes_mean
```

Mean of total day minutes Non Churners: 175.1757543859649

Standard Deviation of total day minutes: 54.45921766392581

Mean of total day minutes for Churners: 206.91407867494823

```
In [23]: 1 #Day Time Minutes Analysis
2
3 day_charge_mean = np.mean(df['total day charge'])
4
5 day_charge_sd = np.std(df['total day charge'])
6
7 non_churn_day_charge_mean = np.mean(df[df['churn'] == 0]['total day cha
8
9 churn_day_charge_mean = np.mean(df[df['churn'] == 1]['total day charge'
10
11 print(f'Mean of total day minutes Non Churners: {non_churn_day_charge_m
12 print(f'Standard Deviation of total day minutes: {day_charge_sd}')
13 print(f'Mean of total day minutes for Churners: {churn_day_charge_mean}
```

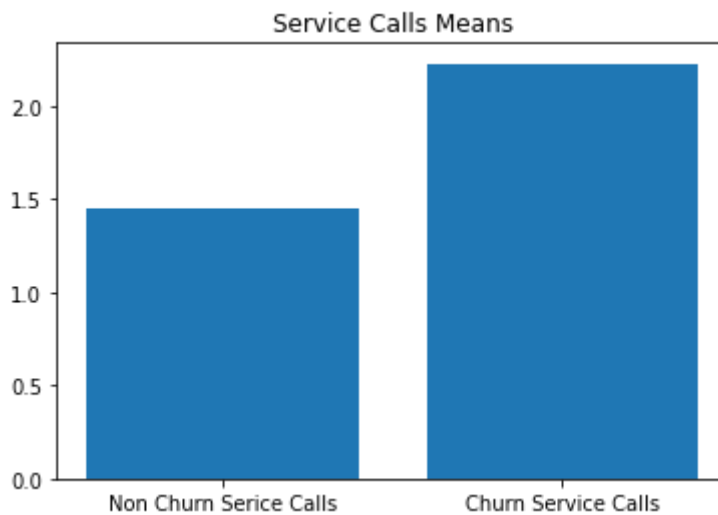
Mean of total day minutes Non Churners: 29.780421052631578

Standard Deviation of total day minutes: 9.258045395636893

Mean of total day minutes for Churners: 35.17592132505176


```
In [24]: 1 import matplotlib.pyplot as plt
2
3 #Service Calls Metrics
4
5 plt.bar(['Non Churn Serice Calls', 'Churn Service Calls'], [
6         non_churn_cust_service_calls_mean, churn_cust_service_calls_mea
7
8 plt.title('Service Calls Means')
```

Out[24]: Text(0.5, 1.0, 'Service Calls Means')



```
In [25]: 1 #Day Minutes Metrics
2
3 plt.bar(['Non Churn Day Minutes', 'Churn Day Minutes'], [
4         non_churn_day_minutes_mean, churn_day_minutes_mean])
5
6 plt.title('Day Minutes Means')
```

Out[25]: Text(0.5, 1.0, 'Day Minutes Means')



## Usage Results

1. Users with higher usage during the day / charges are clearly more inclined to churn
2. This clearly shows that heavy usage users are not experiencing a satisfactory service

## Proposition

1. Provide a service that scales with usage in order to save 15% of the business
2. Specifically when a user starts exceeding 170 / 175 minutes per week offer reduce costs and a higher level of service